

Product Requirements Document (PRD)

Project Title

LLM-Driven Behavioral Segmentation & Product Recommendation Agent for Grupa.io

Objective

Develop an intelligent analytics agent that:

1. Automatically segments users based on behavioral patterns from event logs.
2. Uses an LLM to analyze each cluster and:
 - Identify key pain points
 - Suggest product improvements
3. Presents insights through interactive visualizations.
4. Allows product teams to select a proposed feature to test (via an A/B test planning prompt - no execution yet).

Target Users

- Grupa.io's internal product managers and UX designers
- Growth and strategy teams
- Stakeholders looking to improve user engagement and deal completion

Core Features

1. Data Ingestion:
 - Accepts user logs as .csv: user_id, timestamp, action, metadata
 - Example actions: proposal_create, feature_use, drop_off
2. User Clustering:
 - Applies unsupervised learning (e.g., UMAP + HDBSCAN) to cluster users
 - Feature inputs: session frequency, proposal patterns, feature usage
3. LLM Analysis (Required):
 - For each cluster, send representative behavior summary and metrics
 - LLM (e.g., GPT-4) responds with pain points and product recommendations
4. Interactive Visualization:

- Scatterplot showing user clusters
- Clickable points show insights per cluster

5. A/B Test Selector:

- Select a recommendation to add to the test roadmap (no execution)

Frontend/UX Requirements

- Streamlit or FastAPI + React
- Cluster visualization plot
- Cluster cards with summaries and insights
- Click-to-expand views per cluster
- A/B test planner panel with selection capability

Success Criteria

Clusters generated: ≥ 3 meaningful clusters

LLM insight clarity: pain points + recs per cluster

Visualization: Interactive and explanatory

A/B test planner: Accepts and stores user selections

Tech Stack

- pandas, scikit-learn, umap-learn, hdbscan
- openai, anthropic, langchain
- Streamlit for frontend
- Plotly or Altair for charts
- SessionState or local JSON for test roadmap

Prompt for Cursor

Build a Streamlit app for Grupa.io that does the following:

1. Accepts a CSV with user_id, timestamp, action, metadata.
2. Clusters users using UMAP + HDBSCAN.
3. Sends summaries of each cluster to an LLM for:
 - Pain points
 - Product recommendations
4. Visualizes clusters in a scatterplot.

5. Clicking a cluster shows insights and lets users add recs to an A/B roadmap.
6. No backend execution of A/B test - only planning.
7. Store selections locally. Use placeholder OpenAI key.
8. Use Plotly or Altair and build in modular, readable form.